

Session 6:

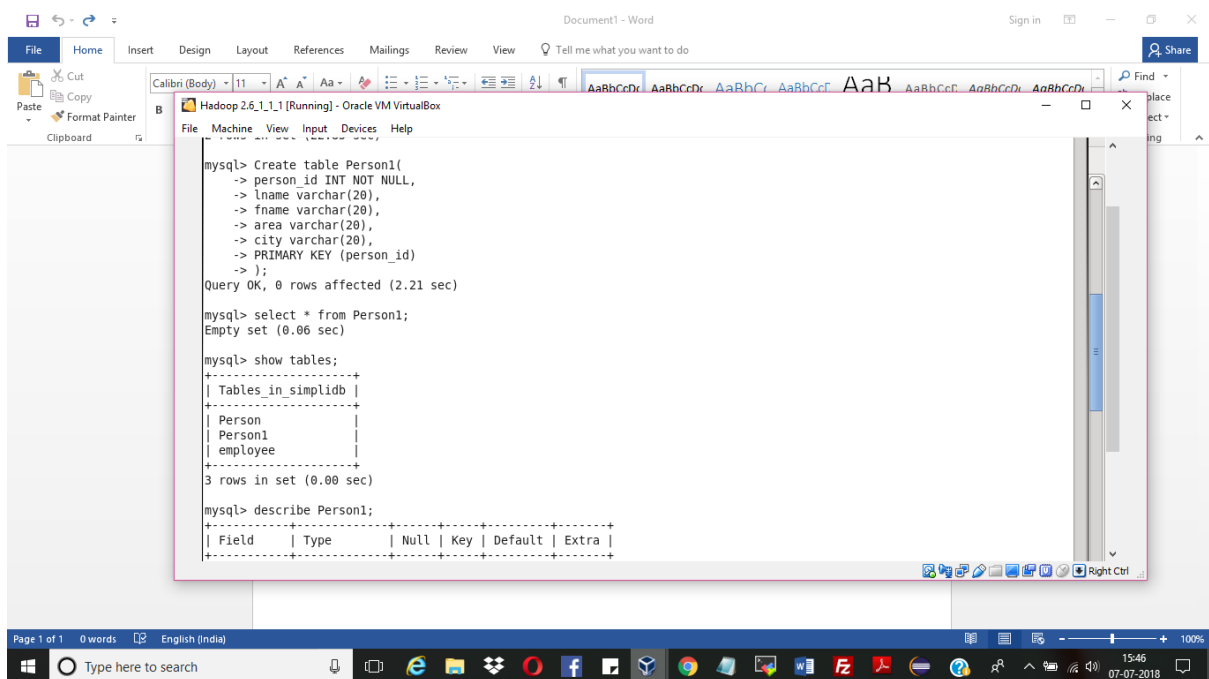
DATA INGESTION TOOL SQOOP & INTRODUCTION TO CASE STUDY

Assignment 1

Task 1

Use Sqoop tool to export data present in SQOOPOUT folder made while demo of Import table

Step 1:



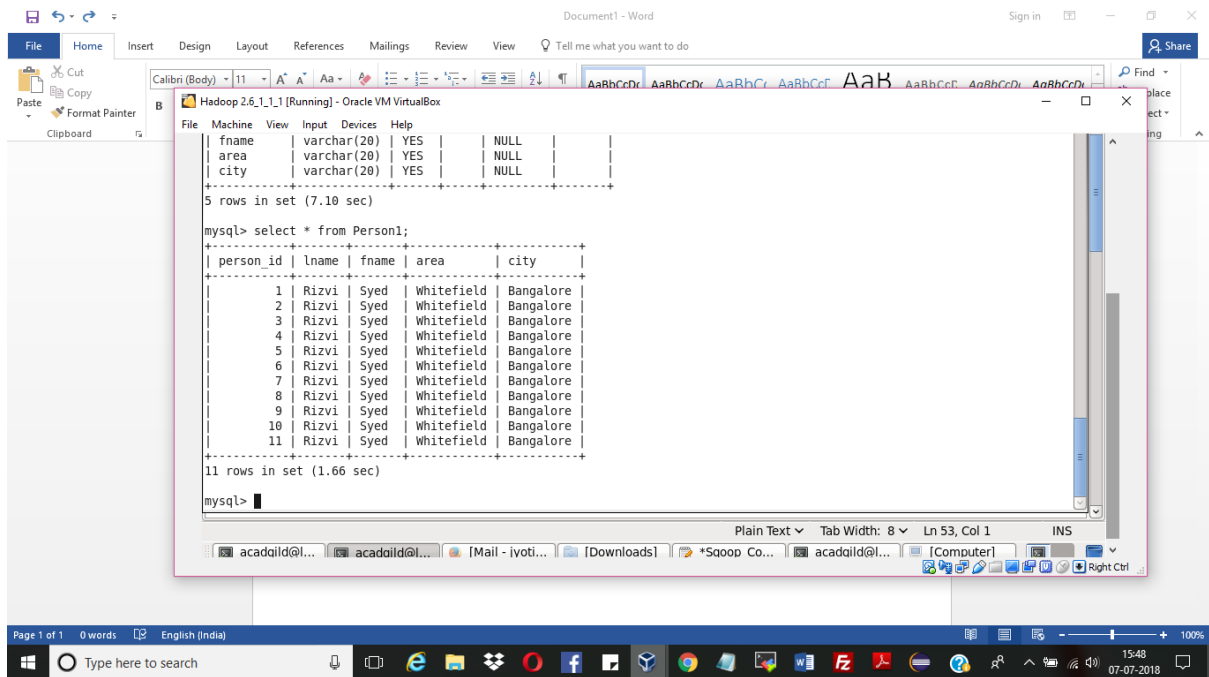
```
mysql> Create table Person1(
-> person_id INT NOT NULL,
-> lname varchar(20),
-> fname varchar(20),
-> area varchar(20),
-> city varchar(20),
-> PRIMARY KEY (person_id)
-> );
Query OK, 0 rows affected (2.21 sec)

mysql> select * from Person1;
Empty set (0.06 sec)

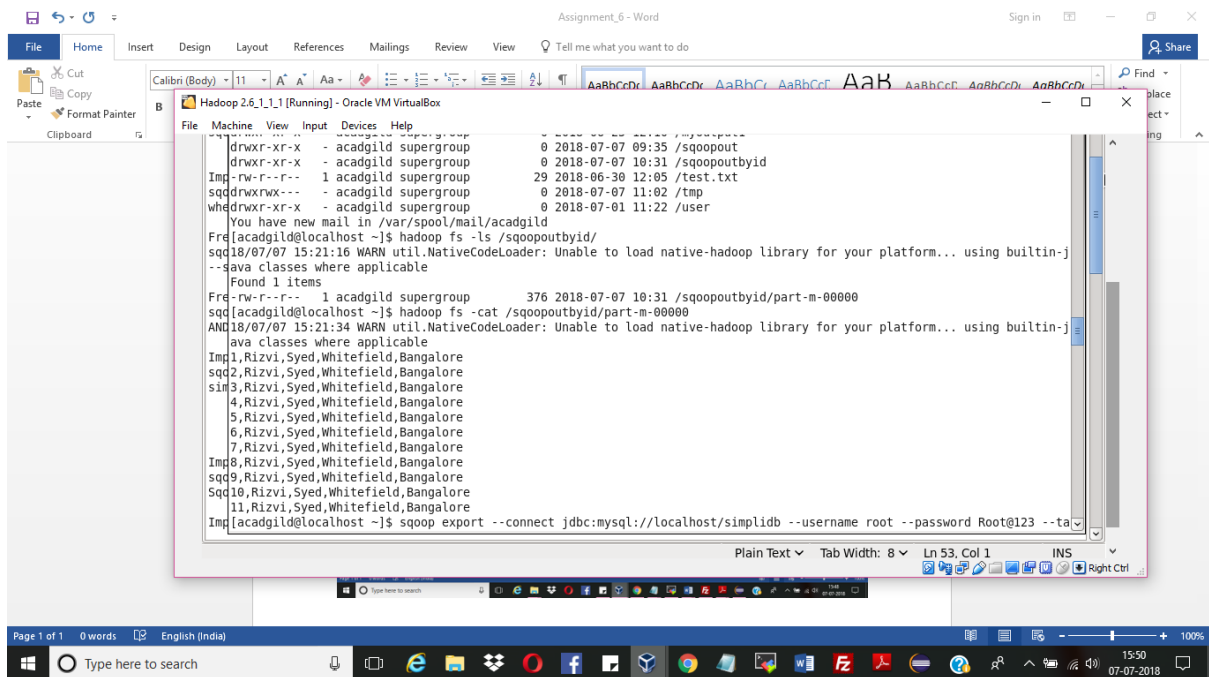
mysql> show tables;
+-----+
| Tables_in_simplidb |
+-----+
| Person              |
| Person1             |
| employee            |
+-----+
3 rows in set (0.00 sec)

mysql> describe Person1;
+-----+
| Field      | Type      | Null | Key | Default | Extra |
+-----+
| person_id  | int(11)   | NO   | PRI |          |       |
| lname      | varchar(20) | YES  |     |          |       |
| fname      | varchar(20) | YES  |     |          |       |
| area       | varchar(20) | YES  |     |          |       |
| city       | varchar(20) | YES  |     |          |       |
+-----+
```

Step 2:



Step 3:



Step 4:


```

18/07/07 15:38:30 INFO mapreduce.JobSubmitter: number of splits:4
18/07/07 15:38:30 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use ma
preduce.map.speculative
18/07/07 15:38:31 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1530934427486_0006
18/07/07 15:38:32 INFO impl.YarnClientImpl: Submitted application application_1530934427486_0006
18/07/07 15:38:32 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1530934427486_0
006/
18/07/07 15:38:32 INFO mapreduce.Job: Running job: job_1530934427486_0006
18/07/07 15:39:18 INFO mapreduce.Job: Job job_1530934427486_0006 running in uber mode : false
18/07/07 15:39:18 INFO mapreduce.Job: map 0% reduce 0%
18/07/07 15:40:30 INFO mapreduce.Job: map 100% reduce 0%
18/07/07 15:40:33 INFO mapreduce.Job: Job job_1530934427486_0006 completed successfully
18/07/07 15:40:34 INFO mapreduce.Job: Counters: 30
    File System Counters
      FILE: Number of bytes read=0
      FILE: Number of bytes written=510244
      FILE: Number of read operations=0
      FILE: Number of large read operations=0
      FILE: Number of write operations=0
      HDFS: Number of bytes read=1464
      HDFS: Number of bytes written=0
      HDFS: Number of read operations=16
      HDFS: Number of large read operations=0
      HDFS: Number of write operations=0
    Job Counters
      Launched map tasks=4
  
```



```

HDFS: Number of bytes read=1464
HDFS: Number of bytes written=0
HDFS: Number of read operations=16
HDFS: Number of large read operations=0
HDFS: Number of write operations=0
Job Counters
  Launched map tasks=4
  Data-local map tasks=4
  Total time spent by all maps in occupied slots (ms)=254461
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=254461
  Total vcore-milliseconds taken by all map tasks=254461
  Total megabyte-milliseconds taken by all map tasks=260568064
Map-Reduce Framework
  Map input records=11
  Map output records=11
  Input split bytes=512
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=3238
  CPU time spent (ms)=9490
  Physical memory (bytes) snapshot=375689216
  Virtual memory (bytes) snapshot=8248029184
  Total committed heap usage (bytes)=130285568
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=0
18/07/07 15:40:34 INFO mapreduce.ExportJobBase: Transferred 1.4297 KB in 131.1166 seconds (11.1656 bytes/sec)
18/07/07 15:40:34 INFO mapreduce.ExportJobBase: Exported 11 records.
You have new mail in /var/spool/mail/acadgild
  
```

Task2

Use Sqoop tool to export data present in SQOOPOUT folder made while demo of Import table with parameter person_id =3.

```

mysql> select * from Person2;
Empty set (0.00 sec)

mysql>
  
```

Check Sqoopoutbyid folder present in hadoop .

```
[acagild@localhost ~]$ hadoop fs -ls /sqoopoutbyid
18/08/19 14:12:40 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 2 items
-rw-r--r-- 1 acagild supergroup 0 2018-08-19 14:12 /sqoopoutbyid/_SUCCESS
-rw-r--r-- 1 acagild supergroup 17 2018-08-19 14:12 /sqoopoutbyid/part-m-00000
[acagild@localhost ~]$ hadoop fs -cat /sqoopoutbyid/part-m-00000
18/08/19 14:13:11 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
USA,NJ,BBB,AAA,3
```

Now we will write export command to fetch data from sqoopoutbyid folder to our mysql PERSON2 table.

```
USA,NJ,BBB,AAA,3
[acagild@localhost ~]$ sqoop export -m 1 --connect jdbc:mysql://localhost/simplidb --username root --P --table Person2 --export-dir /sqoopoutbyid
Warning: /home/acagild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha-../hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /home/acagild/install/sqoop/sqoop-1.4.6.bin_hadoop-2.0.4-alpha-../accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
18/08/19 14:14:02 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6
Enter password:
18/08/19 14:14:07 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.
18/08/19 14:14:07 INFO tool.CodeGenTool: Beginning code generation
Sun Aug 19 14:14:07 IST 2018 WARN: Establishing SSL connection without server's identity verification is not recommended. According to MySQL 5.5.45+, 5.6.26+ and 5.7.6+ requirements SSL connection must be established by default if explicit option isn't set. For compliance with existing applications not using SSL the verifyServerCertificate property is set to 'false'. You need either to explicitly disable SSL by setting useSSL=false, or set useSSL=true and provide truststore for server certificate verification.
18/08/19 14:14:09 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'Person2' AS t LIMIT 1
18/08/19 14:14:09 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM 'Person2' AS t LIMIT 1
18/08/19 14:14:09 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /home/acagild/install/hadoop/hadoop-2.6.5
Note: /tmp/sqoop-acagild/compile/1405dbe0807104bed7c6d0353655b2bf/Person2.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
18/08/19 14:14:12 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-acagild/compile/1405dbe0807104bed7c6d0353655b2bf/Person2.jar
18/08/19 14:14:12 INFO mapreduce.ExportJobBase: Beginning export of Person2
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acagild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acagild/install/hbase/hbase-1.2.6/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
18/08/19 14:14:13 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/08/19 14:14:13 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar
18/08/19 14:14:15 INFO Configuration.deprecation: mapred.reduce.tasks.speculative.execution is deprecated. Instead, use mapreduce.reduce.speculative
18/08/19 14:14:15 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative
18/08/19 14:14:15 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
18/08/19 14:14:15 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/08/19 14:14:15 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps
18/08/19 14:14:15 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/08/19 14:14:17 INFO input.FileInputFormat: Total input paths to process : 1
18/08/19 14:14:17 INFO input.FileInputFormat: Total input paths to process : 1
18/08/19 14:14:18 INFO mapreduce.JobSubmitter: number of splits:1
18/08/19 14:14:18 INFO Configuration.deprecation: mapred.map.tasks.speculative.execution is deprecated. Instead, use mapreduce.map.speculative
18/08/19 14:14:18 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1534667273374_0003
18/08/19 14:14:19 INFO impl.YarnClientImpl: Submitted application application_1534667273374_0003
18/08/19 14:14:19 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1534667273374_0003/
18/08/19 14:14:19 INFO mapreduce.Job: Running job: job_1534667273374_0003
18/08/19 14:14:34 INFO mapreduce.Job: Job job_1534667273374_0003 running in uber mode : false
18/08/19 14:14:34 INFO mapreduce.Job: map 0% reduce 0%
18/08/19 14:14:44 INFO mapreduce.Job: map 100% reduce 0%
18/08/19 14:14:45 INFO mapreduce.Job: Job job_1534667273374_0003 completed successfully
18/08/19 14:14:46 INFO mapreduce.Job: Counters: 30
  File System Counters
    FILE: Number of bytes read=0
    FILE: Number of bytes written=127548
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=148
    HDFS: Number of bytes written=0
    HDFS: Number of read operations=4
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=0
  Job Counters
    Launched map tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=8048
    Total time spent by all reduces in occupied slots (ms)=0
    Total time spent by all map tasks (ms)=8048
    Total vcore-milliseconds taken by all map tasks=8048
    Total megabyte-milliseconds taken by all map tasks=8241152
  Map-Reduce Framework
    Map input records=1
    Map output records=1
```

```

Job Counters
  Launched map tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=8048
  Total time spent by all reduces in occupied slots (ms)=0
  Total time spent by all map tasks (ms)=8048
  Total vcore-milliseconds taken by all map tasks=8048
  Total megabyte-milliseconds taken by all map tasks=8241152
Map-Reduce Framework
  Map input records=1
  Map output records=1
  Input split bytes=128
  Spilled Records=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=124
  CPU time spent (ms)=1950
  Physical memory (bytes) snapshot=105771008
  Virtual memory (bytes) snapshot=2061332480
  Total committed heap usage (bytes)=32571392
File Input Format Counters
  Bytes Read=0
File Output Format Counters
  Bytes Written=0
18/08/19 14:14:46 INFO mapreduce.ExportJobBase: Transferred 148 bytes in 31.0136 seconds (4.7721 bytes/sec)
18/08/19 14:14:46 INFO mapreduce.ExportJobBase: Exported 1 records.
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost ~]$

```

Now Check the Exported data in PERSON2 Table;

```

mysql> select * from Person2;
Empty set (0.00 sec)

mysql> select * from Person2;
+-----+-----+-----+-----+-----+
| person_id | lname | fname | area | city |
+-----+-----+-----+-----+
| 3 | AAA | BBB | USA | NJ |
+-----+-----+-----+-----+
1 row in set (0.00 sec)

mysql>

```